

???2

September 23, 2019

```
[51]: #Exercise 1
import pandas as pd
df = pd.read_csv('https://raw.githubusercontent.com/nickeubank/
→practicaldatascience/master/Example_Data/Datasaurus.txt', delimiter='\t')

[52]: #Exercise2
import numpy as np
for i in range(13):
    x = df['example{}_x'.format(i+1)]
    y = df['example{}_y'.format(i+1)]
    c = pd.concat([x,y],axis=1)
    cr = c.corr()
    print('Example Dataset: {},\nMean x: {},\nMean y: {},\nStd Dev x: {},\nStd_
→Dev y: {},\nCorrelation: {}'.format(i+1,x.mean(),y.mean(),x.std(),y.std(),cr.
→iloc[0,1]))
```

```
Example Dataset: 1,
Mean x: 54.26609978429576,
Mean y: 47.83472062494366,
Std Dev x: 16.769824954043756,
Std Dev y: 26.939743419267103,
Correlation: -0.06412835216739829
Example Dataset: 2,
Mean x: 54.268730022394344,
Mean y: 47.83082315530281,
Std Dev x: 16.7692394934544,
Std Dev y: 26.935726689918788,
Correlation: -0.06858639424107664
Example Dataset: 3,
Mean x: 54.26731970598594,
Mean y: 47.8377172672535,
Std Dev x: 16.760012659806083,
Std Dev y: 26.930036087838204,
Correlation: -0.0683433564802556
Example Dataset: 4,
Mean x: 54.26327323943664,
```


Mean y: 47.832252816901374,  
 Std Dev x: 16.765142039116785,  
 Std Dev y: 26.935403486939123,  
 Correlation: -0.06447185270095167  
 Example Dataset: 5,  
 Mean x: 54.26030345169013,  
 Mean y: 47.83982920901408,  
 Std Dev x: 16.767735488473797,  
 Std Dev y: 26.930191518533462,  
 Correlation: -0.06034144199921762  
 Example Dataset: 6,  
 Mean x: 54.261441783169026,  
 Mean y: 47.83025191366196,  
 Std Dev x: 16.76589790389934,  
 Std Dev y: 26.93987622043797,  
 Correlation: -0.06171483797263012  
 Example Dataset: 7,  
 Mean x: 54.26880527950701,  
 Mean y: 47.83545020401409,  
 Std Dev x: 16.766704015934764,  
 Std Dev y: 26.93999796141102,  
 Correlation: -0.06850422049412323  
 Example Dataset: 8,  
 Mean x: 54.26784882366198,  
 Mean y: 47.83589633112676,  
 Std Dev x: 16.76675894771805,  
 Std Dev y: 26.936104931679974,  
 Correlation: -0.06897973535951203  
 Example Dataset: 9,  
 Mean x: 54.26588178542257,  
 Mean y: 47.83149565232396,  
 Std Dev x: 16.768852670828498,  
 Std Dev y: 26.938608070871844,  
 Correlation: -0.06860920641825635  
 Example Dataset: 10,  
 Mean x: 54.26734110478872,  
 Mean y: 47.83954522535209,  
 Std Dev x: 16.768959216194457,  
 Std Dev y: 26.930274688088435,  
 Correlation: -0.06296110022065422  
 Example Dataset: 11,  
 Mean x: 54.2699272309155,  
 Mean y: 47.836987988408474,  
 Std Dev x: 16.76995861132538,  
 Std Dev y: 26.937683806980512,  
 Correlation: -0.06944556959350363  
 Example Dataset: 12,  
 Mean x: 54.26691630119717,

Mean y: 47.831601987971844,  
 Std Dev x: 16.76999961757302,  
 Std Dev y: 26.937901927731804,  
 Correlation: -0.06657523020460904  
 Example Dataset: 13,  
 Mean x: 54.26015033415493,  
 Mean y: 47.83971727945072,  
 Std Dev x: 16.76995769550748,  
 Std Dev y: 26.930001687162342,  
 Correlation: -0.06558333729297575


Exerciese3 Different datasets have really similar output, which suggests they have highly identical data distribution.

```


[53]: #Exercise4
import warnings
warnings.filterwarnings('ignore', module='plotnine')
from plotnine import *
for i in range(13):
    x = df['example{}_x'.format(i+1)]
    y = df['example{}_y'.format(i+1)]
    p = (ggplot(df, aes(x='example{}_x'.format(i+1), y='example{}_y'.
    →format(i+1)))+
        geom_point())
    print(p)
  
```



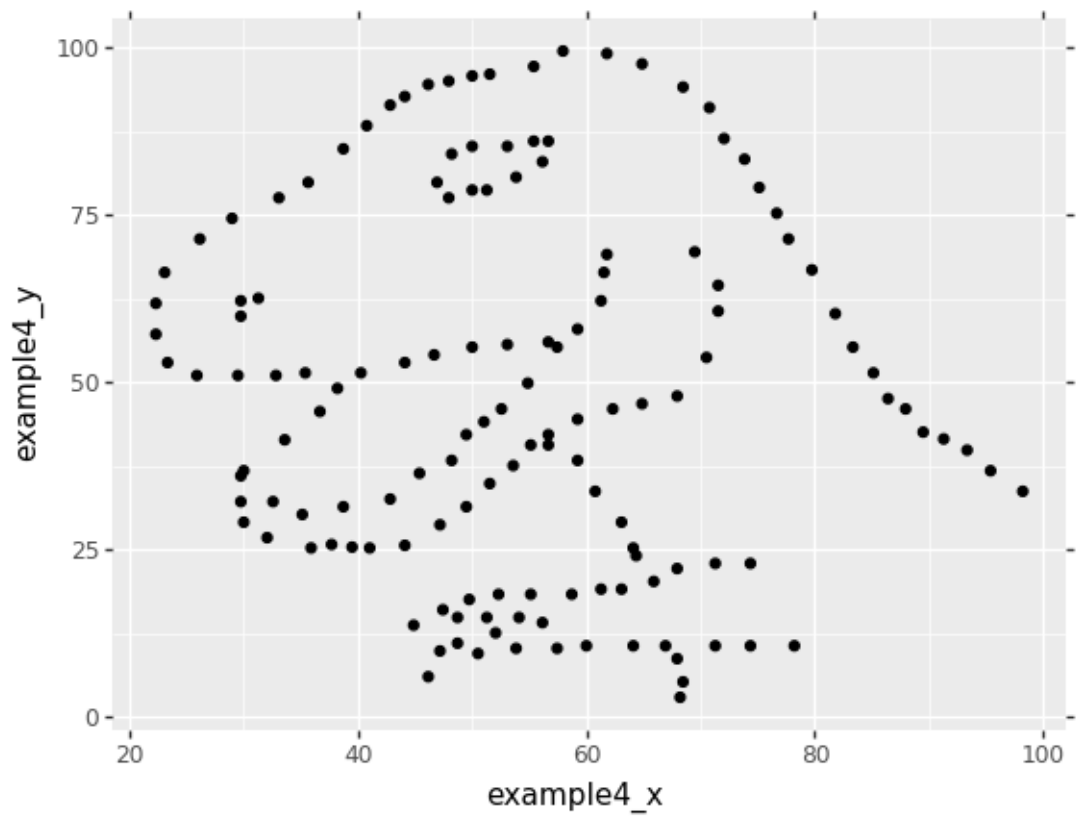
<ggplot: (7551622506)>




<ggplot: (7551071296)>




<ggplot: (7551144403)>



<ggplot: (7552030952)>




<ggplot: (7550969366)>




<ggplot: (-9223372029304070677)>







<ggplot: (7551821730)>




<ggplot: (7551047972)>




<ggplot: (-9223372029304070908)>




<ggplot: (7550958575)>



<ggplot: (-9223372029303431010)>



<ggplot: (-9223372029302988242)>



```
<ggplot: (7551043034)>
```

#Exercise 5 I thought these datasets must look almost the same on the plot. However, these plots show us that data with different distribution can have highly identical statistics.

```
[ ]:
```